

total_frozen_params=210236928
trainable parameters: 28366848
/tmp/ipython-input-32-347286690.py:16: FutureWarning: `torch.cuda.amp.GradScaler(args...)` is deprecated. Please use `torch.amp.GradScaler('cuda', args...)` instead.
self.scaler = GradScaler()
/tmp/ipython-input-32-347286690.py:38: UserWarning: Argument(s) 'always_apply' are not valid for transform Normalize
A.Normalize(mean=[0.5,0.5,0.5],std=[0.5,0.5,0.5],always_apply=True),
DistillationTrainer initialized. Ready to create a distilled dataset.
Loading pre-trained model for distillation...
Creating distilled dataset with top 5.0% hardest examples...

Calculating losses for distillation: 100% 782/782 [06:35<00:00, 2.40it/s]

Original dataset size: 100000
Distilled dataset size: 5000

--- Training on DISTILLED dataset from scratch ---
total_frozen_params=210236928
trainable parameters: 28366848
/tmp/ipython-input-32-347286690.py:16: FutureWarning: `torch.cuda.amp.GradScaler(args...)` is deprecated. Please use `torch.amp.GradScaler('cuda', args...)` instead.
self.scaler = GradScaler()
/tmp/ipython-input-32-347286690.py:38: UserWarning: Argument(s) 'always_apply' are not valid for transform Normalize
A.Normalize(mean=[0.5,0.5,0.5],std=[0.5,0.5,0.5],always_apply=True),
validating: 100% 20/20 [25:00<00:00, 62.77s/it]

train loss: 8.889: 100% 40/40 [00:25<00:00, 2.44it/s]
/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
valid loss: 8.851: 100% 79/79 [00:30<00:00, 3.93it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
0 8.762093 6387.468973 8.731712 6196.328365
saving best model...
train loss: 7.975: 100% 40/40 [00:46<00:00, 3.06it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
valid loss: 7.972: 100% 79/79 [00:30<00:00, 3.78it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
1 8.508478 4956.611654 7.826664 2506.553611
saving best model...
unfreezing GPT2 entirely...
total_frozen_params=0
train loss: 6.058: 100% 40/40 [00:36<00:00, 2.90it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
valid loss: 6.007: 100% 79/79 [00:30<00:00, 3.84it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
2 7.08571 1194.771608 5.698996 298.567547
saving best model...
train loss: 5.890: 100% 40/40 [00:46<00:00, 2.97it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
valid loss: 5.620: 100% 79/79 [00:30<00:00, 3.81it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
3 6.293235 540.900246 5.111638 165.941988
saving best model...
train loss: 5.837: 100% 40/40 [00:46<00:00, 2.90it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
valid loss: 5.123: 100% 79/79 [00:30<00:00, 3.84it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
4 5.759018 317.036972 4.63689 103.222792
saving best model...
train loss: 4.972: 100% 40/40 [00:47<00:00, 2.96it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
valid loss: 4.783: 100% 79/79 [00:30<00:00, 3.80it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
5 5.305374 201.416249 4.341612 76.831289
saving best model...
train loss: 4.766: 100% 40/40 [00:47<00:00, 2.86it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
valid loss: 4.361: 100% 79/79 [00:30<00:00, 3.73it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
6 4.921994 137.276027 3.798373 44.628513
saving best model...
train loss: 4.449: 100% 40/40 [00:47<00:00, 2.80it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():

valid loss: 4.028: 100%

79/79 [00:30<00:00, 3.90it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
7 4.594502 98.938899 3.470978 32.168183
saving best model...

train loss: 4.089: 100%

40/40 [00:46<00:00, 2.93it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():

valid loss: 3.905: 100%

79/79 [00:30<00:00, 3.78it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
8 4.337503 76.516227 3.341515 28.261916
saving best model...

train loss: 3.972: 100%

40/40 [00:46<00:00, 2.99it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():

valid loss: 3.786: 100%

79/79 [00:30<00:00, 3.71it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
9 4.138424 62.703926 3.240253 25.540194
saving best model...

train loss: 4.047: 100%

40/40 [00:47<00:00, 2.82it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():

valid loss: 3.639: 100%

79/79 [00:30<00:00, 3.91it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
10 3.980475 53.542471 3.20237 24.590746
saving best model...

train loss: 3.618: 100%

40/40 [00:46<00:00, 2.95it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():

valid loss: 3.646: 100%

79/79 [00:30<00:00, 3.76it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
11 3.834485 46.269586 3.165132 23.691862
saving best model...

train loss: 4.051: 100%

40/40 [00:46<00:00, 2.88it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():

valid loss: 3.583: 100%

79/79 [00:30<00:00, 4.00it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
12 3.731618 41.74658 3.144382 23.205325
saving best model...

train loss: 3.975: 100%

40/40 [00:47<00:00, 2.91it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():

valid loss: 3.533: 100%

79/79 [00:30<00:00, 3.88it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
13 3.636651 37.964491 3.122258 22.697562
saving best model...

train loss: 3.456: 100%

40/40 [00:46<00:00, 2.92it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():

valid loss: 3.512: 100%

79/79 [00:30<00:00, 3.88it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
14 3.538034 34.399213 3.151534 23.3719
train loss: 3.604: 100%

40/40 [00:25<00:00, 3.02it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():

valid loss: 3.512: 100%

79/79 [00:30<00:00, 3.82it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
15 3.472888 32.229687 3.116772 22.573402
saving best model...

train loss: 3.189: 100%

40/40 [00:46<00:00, 2.88it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():

valid loss: 3.512: 100%

79/79 [00:30<00:00, 3.81it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():
train_loss train_perplexity val_loss val_perplexity
16 3.413554 30.372987 3.127696 22.82135
train loss: 3.288: 100%

40/40 [00:25<00:00, 2.80it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with autocast():

```
valid loss: 3.525: 100% 79/79 [00:29<00:00, 3.81it/s]
/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
  with autocast():
    train_loss train_perplexity val_loss val_perplexity
17 3.385379 29.529189 3.128729 22.844932
train loss: 3.721: 100% 40/40 [00:25<00:00, 2.92it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
  with autocast():

valid loss: 3.537: 100% 79/79 [00:29<00:00, 3.70it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
  with autocast():
    train_loss train_perplexity val_loss val_perplexity
18 3.373603 29.183474 3.134365 22.97405
train loss: 2.947: 100% 40/40 [00:25<00:00, 2.99it/s]

/tmp/ipython-input-32-347286690.py:58: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
  with autocast():

valid loss: 3.530: 100% 79/79 [00:30<00:00, 3.83it/s]

/tmp/ipython-input-32-347286690.py:88: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
  with autocast():
    train_loss train_perplexity val_loss val_perplexity
19 3.348445 28.458456 3.132233 22.925104
Metrics for the model trained on the distilled dataset:
    train_loss train_perplexity val_loss val_perplexity
0 8.762093 6387.468973 8.731712 6196.328365
1 8.508478 4956.611654 7.826664 2506.553611
2 7.08571 1194.771608 5.698996 298.567547
3 6.293235 540.900246 5.111638 165.941988
4 5.759018 317.036972 4.63689 103.222792
5 5.305374 201.416249 4.341612 76.831289
6 4.921994 137.276027 3.798373 44.628513
7 4.594502 98.938899 3.470978 32.168183
8 4.337503 76.516227 3.341515 28.261916
9 4.138424 62.703926 3.240253 25.540194
10 3.980475 53.542471 3.20237 24.590746
11 3.834485 46.269586 3.165132 23.691862
12 3.731618 41.74658 3.144382 23.205325
13 3.636651 37.964491 3.122258 22.697562
14 3.538034 34.399213 3.151534 23.3719
15 3.472888 32.229687 3.116772 22.573402
16 3.413554 30.372987 3.127696 22.82135
17 3.385379 29.529189 3.128729 22.844932
18 3.373603 29.183474 3.134365 22.97405
19 3.348445 28.458456 3.132233 22.925104

--- COMPARISON ---
Full Dataset Training Metrics:
Empty DataFrame
Columns: [train_loss, train_perplexity, val_loss, val_perplexity]
Index: []

Distilled Dataset Training Metrics:
    train_loss train_perplexity val_loss val_perplexity
0 8.762093 6387.468973 8.731712 6196.328365
1 8.508478 4956.611654 7.826664 2506.553611
2 7.08571 1194.771608 5.698996 298.567547
3 6.293235 540.900246 5.111638 165.941988
4 5.759018 317.036972 4.63689 103.222792
5 5.305374 201.416249 4.341612 76.831289
6 4.921994 137.276027 3.798373 44.628513
7 4.594502 98.938899 3.470978 32.168183
8 4.337503 76.516227 3.341515 28.261916
9 4.138424 62.703926 3.240253 25.540194
10 3.980475 53.542471 3.20237 24.590746
11 3.834485 46.269586 3.165132 23.691862
12 3.731618 41.74658 3.144382 23.205325
13 3.636651 37.964491 3.122258 22.697562
14 3.538034 34.399213 3.151534 23.3719
15 3.472888 32.229687 3.116772 22.573402
16 3.413554 30.372987 3.127696 22.82135
17 3.385379 29.529189 3.128729 22.844932
18 3.373603 29.183474 3.134365 22.97405
19 3.348445 28.458456 3.132233 22.925104
```